

# Noise-Busters: How to Sharpen Your Monte Carlo

06009246

Compiled: May 12, 2025

Github Repo: [LINK-TO-THE-TAGGED-RELEASE](#)

## 1 Project Description

This project, **Noise-Busters: How to Sharpen Your Monte Carlo**, offers a hands-on tutorial in R that unpacks three fundamental variance-reduction strategies — antithetic sampling, control variates and importance sampling — using the estimation of the rare-event probability  $P(Z_1 + Z_2 > 4)$ ,  $Z_i \sim \mathcal{N}(0, 1)$  as a running example.

The tutorial consists of a suite of standalone, well-organised scripts. In particular:

1. `simulations.R`: defines the baseline sampler and auxiliary simulators;
2. `variance_reduction_methods.R`: implements each reduction algorithm;
3. `confidence_intervals.R`: constructs data frames of estimates with confidence bounds;
4. `plotting.R`: produces clear, annotated ggplot2 visualisations;
5. `true_value.R`: supplies the analytical probability for reference.

Each script is modular, immediately executable and fully documented with roxygen2-style comments. The tutorial itself is authored as a Quarto notebook, weaving explanatory text, reproducible code and publication-quality plots to demonstrate the  $O(1/\sqrt{n})$  convergence bottleneck of naïve Monte Carlo and to compare how each technique accelerates convergence and curtails sampling variability.

**A strong emphasis is placed on reproducibility and sound software engineering: simulations are seeded for exact replication, computationally intensive chunks are cached, and the environment is locked down with `renv`.** Rather than embedding the R functions implementing each variance-reduction technique directly in the narrative, the tutorial focuses on conceptual exposition and visual demonstration; interested readers are most welcome to explore the complete codebase and accompanying materials in the project's GitHub repository.

## 2 Assessment Criteria

### Technical Competence:

- Developed a modular R codebase (in `R/`) implementing standard Monte Carlo and three variance-reduction methods.
- Instrumented convergence diagnostics via `confidence_intervals.R` to compute and tabulate mean  $\pm$  CI over increasing sample sizes.

### User Interface:

- Authored a Quarto notebook with a clear, engaging structure.
- Encouraged curious readers to inspect every standalone R script via the GitHub repository, fostering hands-on exploration of the utilised functions.

### Analysis and Interpretation:

- Framed each section with precise questions, contrasted empirical estimates with the closed-form true value (`true_estimate`), and interpreted CI widths in relative terms.
- Compared bias, variance, and compute time across methods — to highlight not only the variance precision improvement of the discussed methods but also their computational efficiency.

### Presentation and Communication:

- Followed a narrative arc (“Hook  $\rightarrow$  Chaos  $\rightarrow$  Methods  $\rightarrow$  Comparison  $\rightarrow$  Takeaways”) with short text blocks, call-outs, and transition headings to maintain reader engagement.
- Produced simple but informative ggplot2 visualisations in `plotting.R`, annotated with true-value lines and concise axis labels, ensuring immediate interpretability.

### Reproducibility and Documentation:

- Managed package versions with `renv`, seeded all random draws for exact replication, and cached expensive chunks to guarantee consistent outputs across renders.
- Documented every function with roxygen2-style comments (in `R/*.R`), and provided a one-click Quickstart in `README.md` to restore the environment and render the notebook.

### Project Management:

- Organised the repository into logical folders; create and followed a milestone plan with concrete deadlines.
- Captured the finished tutorial in a single comprehensive commit and marked the submission milestone with a Git tag: `v1.0.1`.

### 3 Project Reflection

#### Learnings:

- I discovered how orchestrating modular scripts—from samplers to plotting — creates a flexible framework that can be extended to new Monte Carlo problems with minimal friction.
- Creating convergence diagnostics and CI visualisations reinforced the practical importance of tracking error bounds as a function of  $n$ , rather than relying on point estimates alone.
- Embedding roxygen2 documentation and renv-driven reproducibility early on made the codebase feel like a polished software library rather than a one-off analysis notebook.

#### Challenges:

- Initially, I scoped the tutorial as three parts—introducing simple Monte Carlo via the  $\pi$  example, detailing variance-reduction techniques, and applying them to a real-world problem — but quickly realised this plan was too ambitious and refocused on the core rare-event narrative for depth over breadth.
- I did not allocate any time to writing `testthat` tests, leaving key simulation functions unverified; I now appreciate that formal tests are crucial for catching regressions and ensuring statistical correctness, and would include comprehensive test coverage from day one in any future project.

#### Further Development:

- Going forward, I am keen to explore interactive Shiny or Voila widgets that let users experiment with sample sizes and variance-reduction parameters in real time, turning static plots into dynamic learning tools.